

# Computer Graphics

## Lecture 4

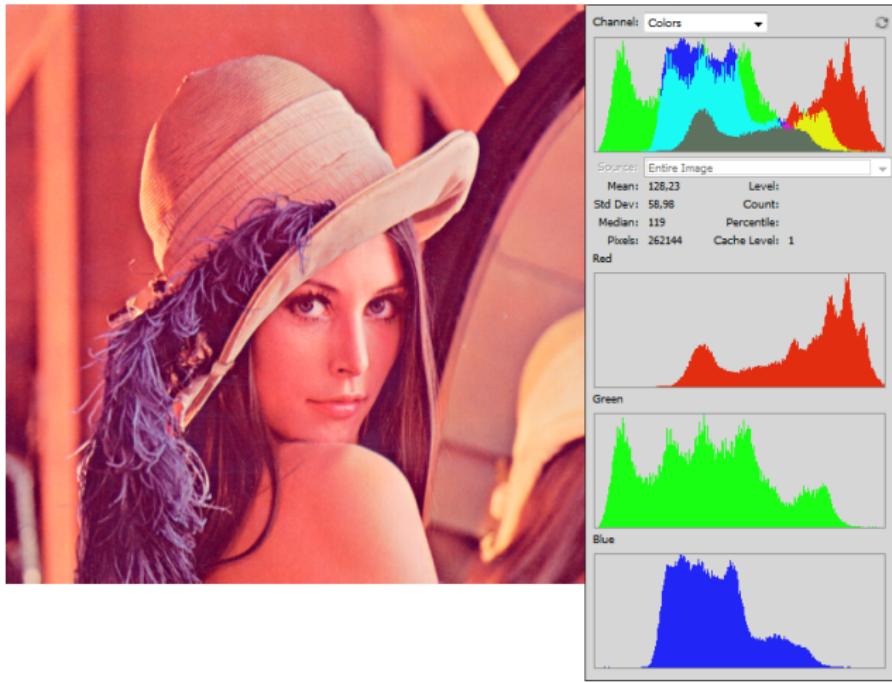
Paweł Aszklar  
P.Aszklar@mini.pw.edu.pl

Faculty of Mathematics and Information Science  
Warsaw University of Technology

Warsaw 2013

# Histogram

- Graphical representation of tonal distribution
- Plots the number of pixels in each tonal value



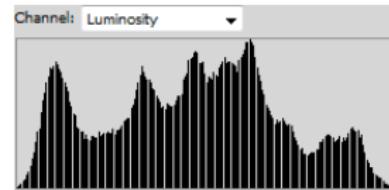
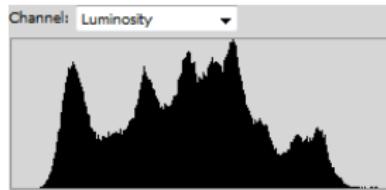
# Histogram Stretching

- Equivalent as contrast enhancement filter
- All channels stretched equally or each channel stretched independently
- Image transformed to use the full range of available intensity values

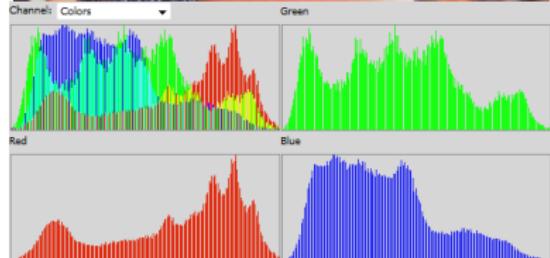
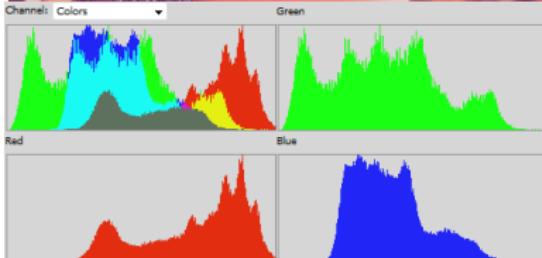
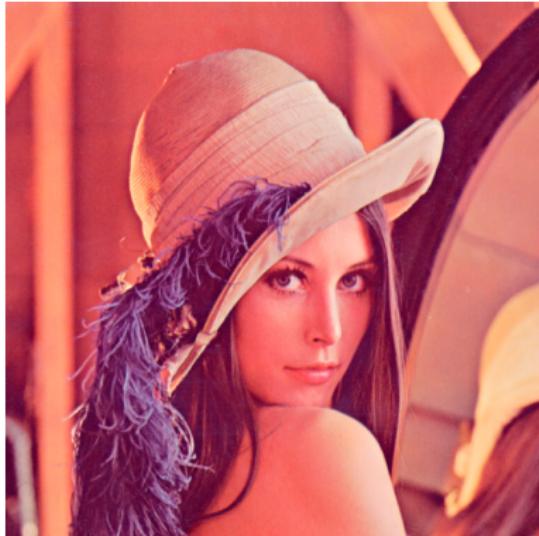
$$I(x, y) = \frac{255}{I_{\max} - I_{\min}} (I(x, y) - I_{\min})$$

- When determining values of  $I_{\min}$  and  $I_{\max}$  we can ignore intensities for which the histogram value is lower than a certain threshold

# Histogram Stretching



# Histogram Stretching



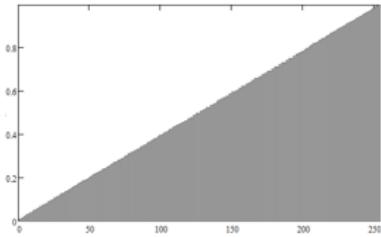
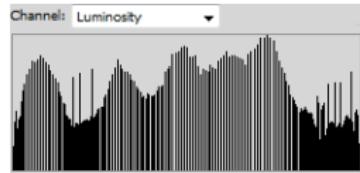
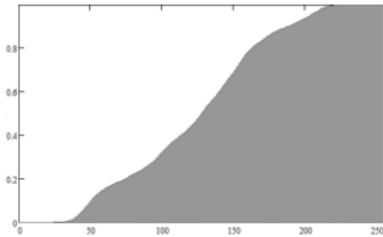
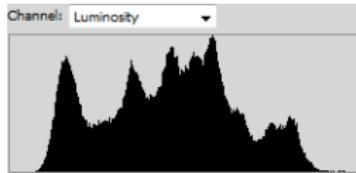
# Histogram Equalization

- The goal is to get more uniform distribution of intensities across the available range
- Effectively spreads out most frequent intensity values on the histogram
- Increases the contrast of lower local contrast
- Improves the linearity of the cumulative distribution function

$$D(i) = \frac{\sum h(i)}{n_{\text{pixels}}}$$

$$I(i) = \frac{D(i) - D_0}{1 - D_0} (K - 1)$$

# Histogram Equalization



# Histogram Equalization



# Color Quantization

- Process of representing an image with a limited number of colors
- The set of colors used is called a *color palette*
- Applications:
  - Displaying images on (legacy or industrial) hardware supporting restricted number of colors (often due to memory and processing power limitations rather than display capabilities)
  - Conversion to file formats supporting color palettes
  - Image and video compression
- First step is to select the color palette
- Second step is to replace each pixel with a color from the palette
  - Usually nearest in palette neighbor search (any norm, e.g.: Euclidean norm, Manhattan norm, etc.) in any color space (usually RGB)
  - Euclidean distance in CIELAB color space often gives the best results

# Color Quantization Methods

## Fixed Quantization

- A predefined color palette (imposed by hardware or chosen arbitrarily) and fixed color mapping
- Potential artifacts: color shift, false contour, etc.
- E.g.: Uniform quantization

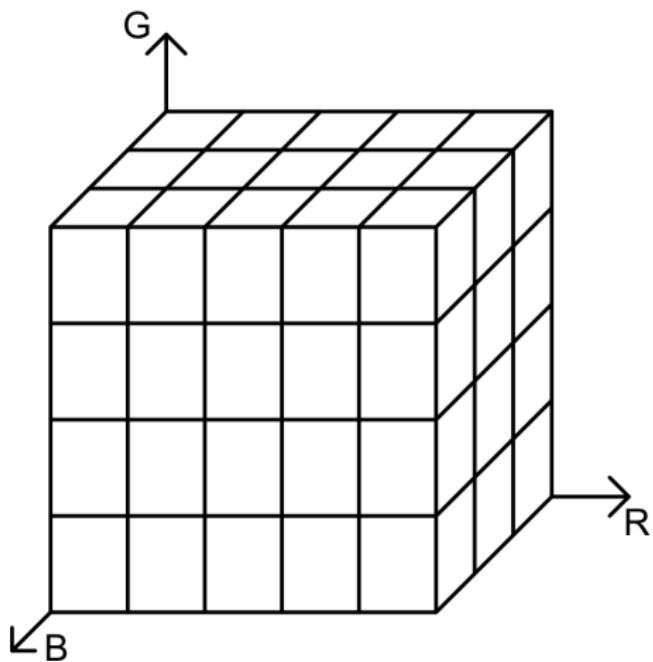
# Color Quantization Methods

## Adaptive Quantization

- Selects image-dependent color palette
- Better quality, but more time consuming
- Any three-dimensional clustering algorithm will work (octree, popularity algorithm, median-cut, k-means)
- Specialized algorithms:
  - Spatial Color Quantization <http://www.cs.berkeley.edu/~dcoetzee/downloads/scolorq/>
  - NeuQuant Algorithm  
<http://members.ozemail.com.au/~dekker/NEUQUANT.HTML>

# Uniform Quantization

- Arbitrary palette, image-independent
- Uniform subdivision of RGB cube for each coordinate
- Number of colors:  $k_R \times k_G \times k_B$
- Palette colors — centers of cuboids



# Popularity Algorithm

- Adaptive method
- Palette consists of the K most frequently occurring colors
  - To improve performance only check colors at local maxima of the histograms
  - Alternatively perform preliminary uniform subdivision (like uniform quantization) and search for cuboids with the most points in them
- Data structures:
  - Linked list — very high time complexity
  - Array indexed with colors — very high space complexity
  - BST, Octree

# Octree Color Quantization

- Building Octree of colors

- Colors of pixels are sequentially added to the leaves of the octree
- If no leaf is found for a color, a new leaf is inserted at certain depth (usually depth of 8 for 8-bit per channel RGB)
- Leaves store the count of pixels added to them and the sum of their colors
- If after an insertion the number of leaves exceeds the size of the palette, reduce the tree until it doesn't

```
void insert(OctreeNode** node, Color c, int depth)
{
    if (*node == null)
        CreateAndInit(node, depth);
    if (isLeaf(*node)) {
        (*node)->pixels_count++;
        addColors(*node, c);
    } else
        insert(&(*node)->child[Branch(c, depth)], c, depth + 1);
}
```

# Octree Color Quantization — Tree Reduction

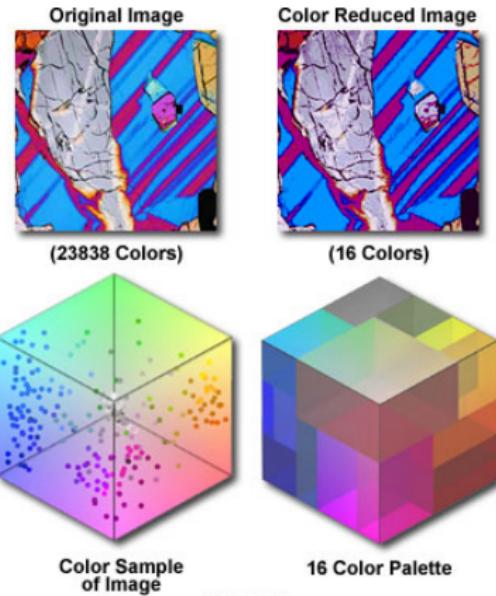
- Octree reduction is performed if number of leaves exceeds the available size of the color palette
- Reduction criteria
  - Select a non-leaf node with the highest depth value
  - If more than one exists, chose the one representing the fewest pixels (or the most, or any)
  - Make that node a leaf (number of pixels and the sum of colors is the total of values from children)
- Reduction may need to be repeated if selected node had only one child

# Octree Color Quantization — Color Mapping

- After the tree has been built, for each leaf select a color in the final palette (average of colors added to that leaf)
- Color mapping — for each pixel traverse the final tree again to find a corresponding leaf. Assign to that pixel a color from the palette represented by that leaf

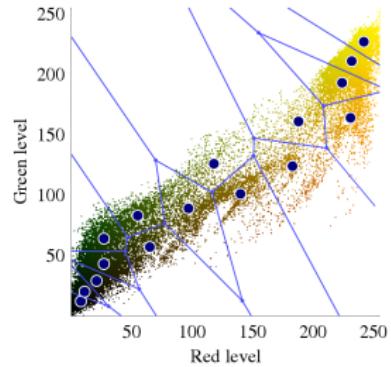
# Median-Cut Algorithm

- Subdivide the color space (e.g. the RGB cube) into  $K$  cuboids in such a way that each of them contains (approximately) the same number of pixels
- Each cuboid represents a color in the palette (either center of the cuboid or average of pixels)



# K-Means Algorithm

- Well-known statistical algorithm
- Iterative
- Select  $K$  random colors — *centroids* ( $K$  — size of the color palette)
- Assign each pixel to its closest centroid
- For each group calculate the center of gravity (average) which becomes a new centroid (replacing the old one)
- Reassign the pixels to centroids and repeat until no changes occur



# Color Quantization and Dithering

- Once the color palette is established, we can always use some metric (e.g. Euclidean distance) to select a color from the palette for each pixel
- This makes it easy to incorporate dithering into the color selection (especially using any of the error propagation methods)

# Color Quantization and Dithering

Example:

- Color palette size = 8
- 2 levels per channel

Original image:



Octree quantization:



Error diffusion:



# Color Quantization and Dithering

Example:

- Color palette size = 256
- Levels per channel — 6R/7G/6B

Original image:



Octree quantization:



Error diffusion:

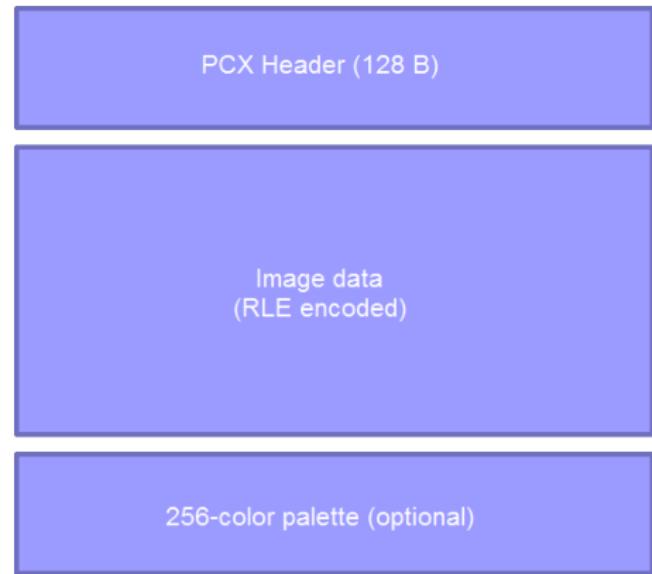


# Popular Raster Image File Formats

- PCX (legacy)
- BMP
- GIF
- PNG
- JPEG

# PCX

- Personal Computer eXchange
- Developed by ZSoft company
- Simple, fast, lossless compression
- Designed for 16 (stored in header) or 256 (appended at the end) color palette, extended to true-color (24-bit) images
- Image data compressed using run-length encoding algorithm (RLE)



# PCX — RLE Compression

- Bytes represent either
  - Value of the pixel (or palette index)
  - Beginning of a pair. First byte is the count, second is the value
- First two bits set indicate a pair, for example (240, 167) indicate that value 167 should be repeated  $240 - 191 = 51$  times.
- Some values (greater than 190) need to be encoded as a pair even if they are not repeating
- Allows for 0 run-length data

# Bitmap (BMP)

- File structure
  - File header
  - Bitmap information header
  - Optional palette
  - Image data
  - Some versions allow to embed ICC color profile
- Huge file of uncompressed images
- Lossless — no quality loss during multiple edits, however other lossless formats are preferred nowadays
- Color components stored in reverse order:  $[B, G, R]$
- Row of data always stored as a multiple of 4 bytes

# Bitmap Header

```
typedef struct {
    unsigned short int type;          /* Signature (usually BM) */
    unsigned int size;                /* File size in bytes */
    unsigned short int reserved1, reserved2;
    unsigned int offset;              /* Offset to image data in bytes */
} HEADER;

typedef struct {
    unsigned int size;                /* Header size in bytes */
    int width, height;               /* Image width and height */
    unsigned short int planes;        /* Number of color planes (usually 1) */
    unsigned short int bits;          /* Bits per pixel */
    unsigned int compression;         /* Compression type */
    unsigned int imagesize;           /* Image size in bytes */
    int xresolution, yresolution;    /* Pixels per meter */
    unsigned int ncolours;            /* Number of colors in palette */
    unsigned int importantcolours;    /* Number of significant colors in palette */
} INFOHEADER;
```

# GIF

- **G**raphics **I**nterchange **F**ormat
- Developed in 1987 by CompuServe company
- Uses 256 global color palette
- Allows to store multiple images in blocks
- Each block can have it's own local 256-color palette
- Allows for extension blocks supporting animations, selecting one transparent color, interlacing etc.
- Uses LZW (Lempel, Ziv, Welch, '84) compression — it was subject of patent claims by Unisys company, but as of 2004 those patents expired.
- LZW is a lossless dictionary compression algorithm
- Good for graphics, logos, etc. Poor for photographs (difficult to represent true-color images)

# GIF - LZW compression

Encoding:

- ① Initialize dictionary with data alphabet

- ②  $q = \text{first input symbol}$

- ③ While input not empty

- ① read symbol  $s$

- ② if  $q+s$  in dictionary,

$$q=q+s$$

- ③ otherwise

- ④ output dictionary code for  $q$
    - ⑤ add  $q+s$  to dictionary
    - ⑥  $q=s$

- ④ output dictionary code for  $q$

Decoding:

- ① Initialize dictionary with data alphabet

- ②  $pc = \text{first input code}$

- ③ Output dictionary entry for  $pc$

- ④ While input not empty

- ① read code  $c$

$$pq=\text{dictionary}[pc]$$

- ② if code  $c$  in dictionary, add ( $pq+\text{first symbol of dictionary}[c]$ ) to dictionary and output  $\text{dictionary}[c]$

- ③ otherwise add ( $pq+\text{first symbol of } pq$ ) to dictionary (special case)

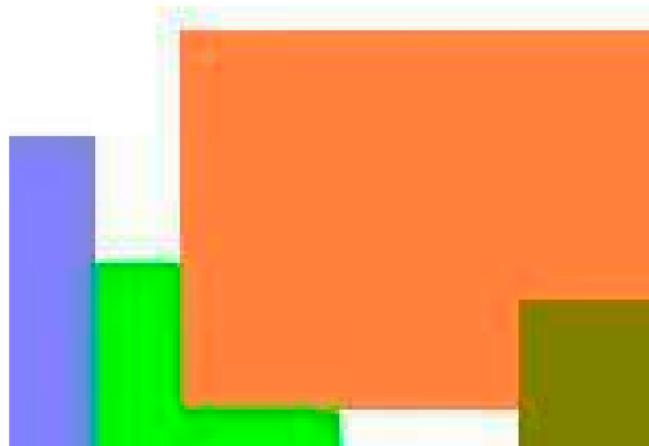
$$pc=c$$

# PNG

- Portable Network Graphics
- Developed as an alternative to GIF when it was still subject to patents
- Two step compression algorithm
  - pre-compression filtering (usually improves compression)
  - dictionary compression using DEFLATE algorithm (modified LZ77)
- Supports grayscale and RGB images (up to 16-bit per channel; both w/ and w/o alpha or with single transparent color) and color palettes (with optional transparency settings for colors in palette)
- Improved interlacing (over GIF), no animations
- Allows for extensions specifying absolute color space and white point, gamma correction, histograms, color profiles, pixel aspect ratio, stereoscopy and many more features

# JPEG

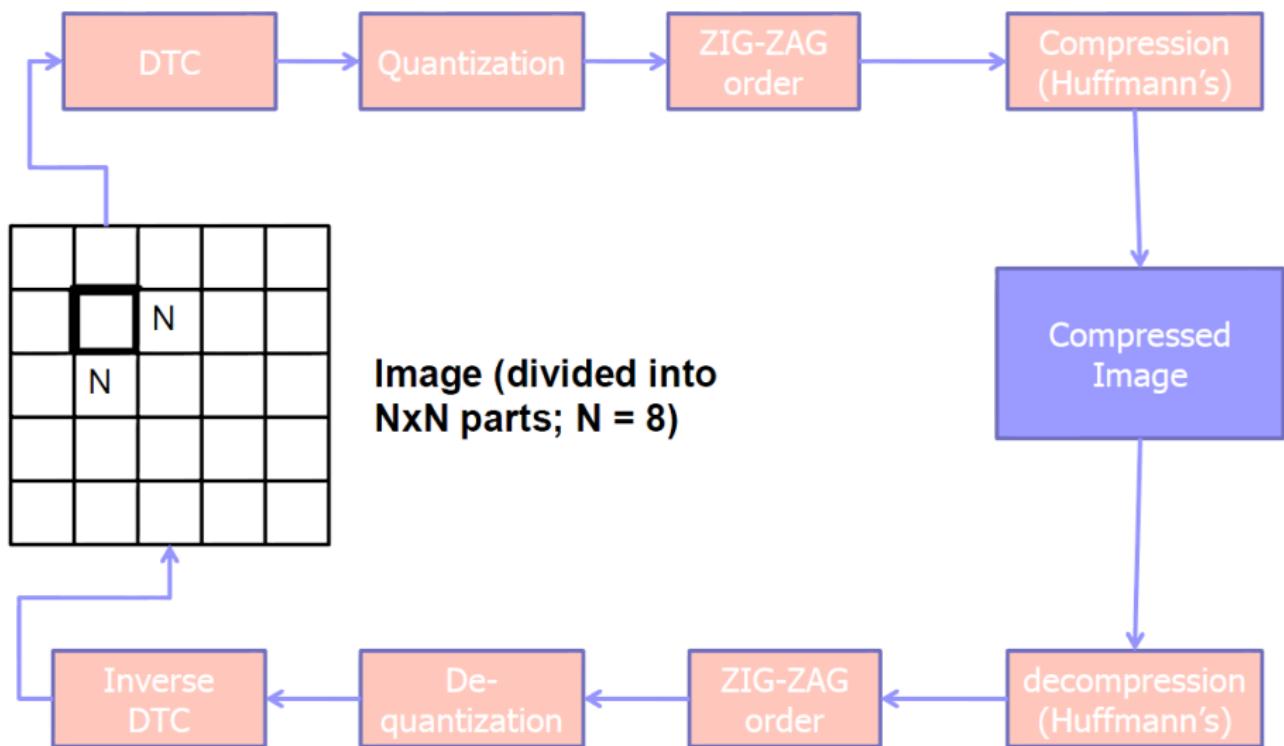
- Lossy compression method for digital photography created by **Joint Photography Experts Group**
- Very good compression with acceptable quality loss for photography. Poor performance with more regular graphics, logos, etc.



# JPEG Compression Algorithm

- ① Image data converted into luminance and two chromaticity components ( $Y'CrCb$ )
- ② Image channels is divided into  $N \times N$  blocks
- ③ Discrete Cosine Transform (DCT) is applied to every block — Instead of pixel values we receive value means and frequencies of value variations within the blocks. No data loss here.
- ④ Data is subject to quantization — floating-point values are replaced with integers. **Here is where the data loss takes place.**
- ⑤ Zig-zag order is applied to DTC coefficients, so 0 values lie close together.
- ⑥ Non-zero coefficients are encoded (losslessly) using Huffmann's coding

# JPEG Compression/Decompression



End of lecture 4